



# Chapter 1

## History and Building Blocks of Cloud Computing

Mastering Cloud Computing

Coleman Kane

(based on materials from Paul Talaga)


# 1. Distributed Systems

*A distributed system is a collection of independent computers that appears to its users as a single coherent system. - Tanenbaum*

To share resources for better utilization.

We have a course in Distributed Systems!

# What is old is new again...



# Mainframes (1951 onward)

- Multiple processing units
- Powerful, reliable, IO optimized
- Timesharing systems
- Replacement of components while running - always on
- Still used for transaction processing: banking, airline ticketing, registration

IBM z990  
2003  
256 GB RAM  
1k+ VM Linux  
11k SSL con/sec



IBM EC12  
2012  
101 CPUs  
5.5Ghz hex core  
Integrated SSDs  
6k Linux VMs

# Clusters (1968 onwards)

- Networked commodity (cheap) machines
- Physically close (room/building/LAN)
- Good for processing work, not IO
- Easy to expand
- Some use spare cpu cycles
- [Condor](#), [Beowulf Clusters](#), Message Passing Interface [\(MPI\)](#)
- [Ohio Supercomputer](#)

# Grid Computing (1990s onwards)

- Like cluster, but heterogeneous nodes
- Large physical distances
- Utility computing idea
- Needed high bandwidth connections (Internet)
- Good for processing work, not IO
- [SETI@Home](#), [BOINC](#)

# But what about latency!?!

Talaga's [Dissertation](#)/Digression

- A semi full of hard drives has HUGE bandwidth!
- We get sold with bandwidth, but low latency is what we're after.
- But the speed of light is fixed!
  - Significant drop in sales for > 500ms response
  - Google goal is < 200ms
  - NY to LA is 74ms RTT
- Grids may be good for some workloads, but HORRIBLE for others (IO and communication)
- Mainframe->Cluster->Grid, latency between nodes increases.
- Data location matters, and how your app uses data.

# Cloud Computing

- Next evolution after grid computing
- Infinite capacity
- Resilient to failures
- Always on
- Built using commodity machines
- Pay-per-use (Utility vision)



# True root (my view)

The implementation and wider use of **distributed systems** theory to build distributed databases and file systems allowed cloud computing to take off.

- Huge datasets (multiple petabytes)
- Vector clocks ([Lamport timestamps](#))
- Paxos ([Wiki](#))

## 2. Virtualization

- Abstract core computing elements away
  - Processor
  - Storage
  - Networking
- Hardware virtualization (VMware, VirtualBox, XEN, EC2, etc....)
  - Most performance issues solved
- Process virtualization (Google AppEngine, Azure, Java)

# 3. Web 2.0

- Web 1.0? - Static pages
- Web 2.0 is:
  - “Web as platform” - *John Battelle and Tim O'Reilly*
  - Interactivity & flexibility - Allow users to change a site's content!
  - Asynchronous JavaScript and XML (AJAX)
  - Web Services
- Not a ‘next version’ of the web, but a way of using HTTP/HTML
- Examples: Google Docs, Facebook, YouTube, Wikipedia

# 4. Service-oriented Computing

A component that can perform any function.

- Loosely coupled
- Reusable
- Programming language independent
- Location transparent

By layering services we can build a service-oriented architecture (SOA)

# Service-Oriented Computing

Important attributes:

- Quality of Service (QoS)
  - Service attributes, response times, security, uptime, etc...
- New software delivery model
  - Can sell components, not entire programs
  - Access through the internet
  - HTTP
  - Web Service Description Language (WSDL)
  - Simple Object Access Protocol (SOAP)

# 5. Utility-oriented Computing

Storage, compute, applications, infrastructure, all on a pay-per-use basis.

Old idea.. job queueing systems, OS time-slicing, all developed in mainframes to charge for use.

Buying services/products online common now.

# Current Platforms

- Amazon Web Services (AWS) - Current leader in IaaS: Elastic Compute Cloud (EC2), Simple Storage Service (S3), many many others.....
- Google AppEngine - PaaS - Python, Java, Go
- MS Azure - IaaS & PaaS
- Hadoop - Open Source framework for MapReduce
- Force.com & Salesforce.com

# Serverless Architectures

- *Microservice architecture*
- CloudFoundry
- Amazon Lambda
- BlueMix
- Predix
- Microsoft Services Framework



# Some Security Implications

- Provisioning in software - automated, online, hands-off
- Full-system imaging & snapshotting, assists forensics investigation
- Partitioning - Better isolation between applications than virtual web hosting servers
- Software networking - pros / cons