

Chapter 3

VMM, OS Level Virtualization, and Summary


Mastering Cloud Computing
Coleman Kane

(based on material by Paul Talaga)

VMM in Detail

3 parts:

- Dispatcher (route instructions)
- Allocator
 - system resource Ok?
- Interpreter routines
 - What to do with sensitive instructions? - run special code!




VMM Requirements - by Goldberg and Popek (1974)

- *Equivalence* - A guest should behave exactly as it did on physical hardware
- *Resource control* - VMM should be in complete control
- *Efficiency* - Significant amount of guest instructions untouched by VMM

Goldberg and Popek's 3 Theorems for a VMM

1. *For any conventional third-generation computer (IC), a VMM may be constructed if the set of sensitive instructions for that computer is a subset of the set of privileged instructions.*



Goldberg and Popek's 3 Theorems for a VMM

2. A conventional third-generation computer is recursively virtualizable if:

1. It is virtualizable
2. A VMM without any timing dependencies can be constructed for it.

A VM in a VM (in a VM)? YES! But will be slow.

Goldberg and Popek's 3 Theorems for a VMM

3. A hybrid VMM may be constructed for any conventional third-generation machine in which the set of user-sensitive instructions is a subset of the set of privileged instructions.

A hybrid virtual machine (HVM) traps user-level instructions as well.

Hardware Virtualization Techniques

- Hardware Assisted
- Full virtualization
- Paravirtualization
- Partial virtualization

Hardware Assisted

- Hardware provides features for virtualization - Intel VT & AMD V
- Adds extensions to x86-64
- Originally used in [IBM System/370](#)
- Examples: Kernel-based VM (KVM), VirtualBox, Xen, VMware, Hyper-V, Sun xVM, Parallels

Full Virtualization

- Ability to run guest with no modifications.
- Guest has no idea it is virtualized.
- All privileged instructions need to be trapped
- Required for Windows OSs when HW not available.
- Works on all hardware, but SLOW

Paravirtualization

- Non-transparent virtualization
- Guest must be changed - change sensitive instructions to VMM API calls
- Simpler VMM
- Much faster:
 - No traps! Code runs on bare hardware
 - OS/driver optimizations for greater speed
- Installing guest drivers makes full virtualization into para-ish
- Examples: Xen, IBM VM OS, VMware, Parallels

Partial Virtualization


- Not all features of hardware/OS are virtualized
- Examples:
 - Address space virtualization - virtual mem
 - RAID - disk
 - Network bonding
 - VPNs
 - Same hardware otherwise

Practical Issues


- Hardware & Full Virtualization
 - Can overprovision CPU
 - Dedicated Ram
 - Dedicated disk (sometimes, delta drives possible)
- Paravirtualization
 - Can overprovision CPU & **RAM!**
 - Can set RAM ranges and let OSs self manage!

OS-level Virtualization

- Separate user-spaces
- No hypervisor needed
- Different filesystems, network configurations, software, access to devices
- Extension of the [chroot](#) unix concept
- **Little to no overhead**
- Can use any hardware, any software
- But, must all use same OS
- Ex: FreeBSD Jails, Parallels Virtuozzo Containers, OpenVZ, EC2, Docker, [others..](#)




- Open Sourced Mar 2013 - Jan '14 capital
- Extension of LXC (Linux Containers)
- Uses Images to build other Images (Containers)
- Git feel
- Simple config file to generate images
- Lightweight & processes run on bare hardware (fast!)



Virtual Machines

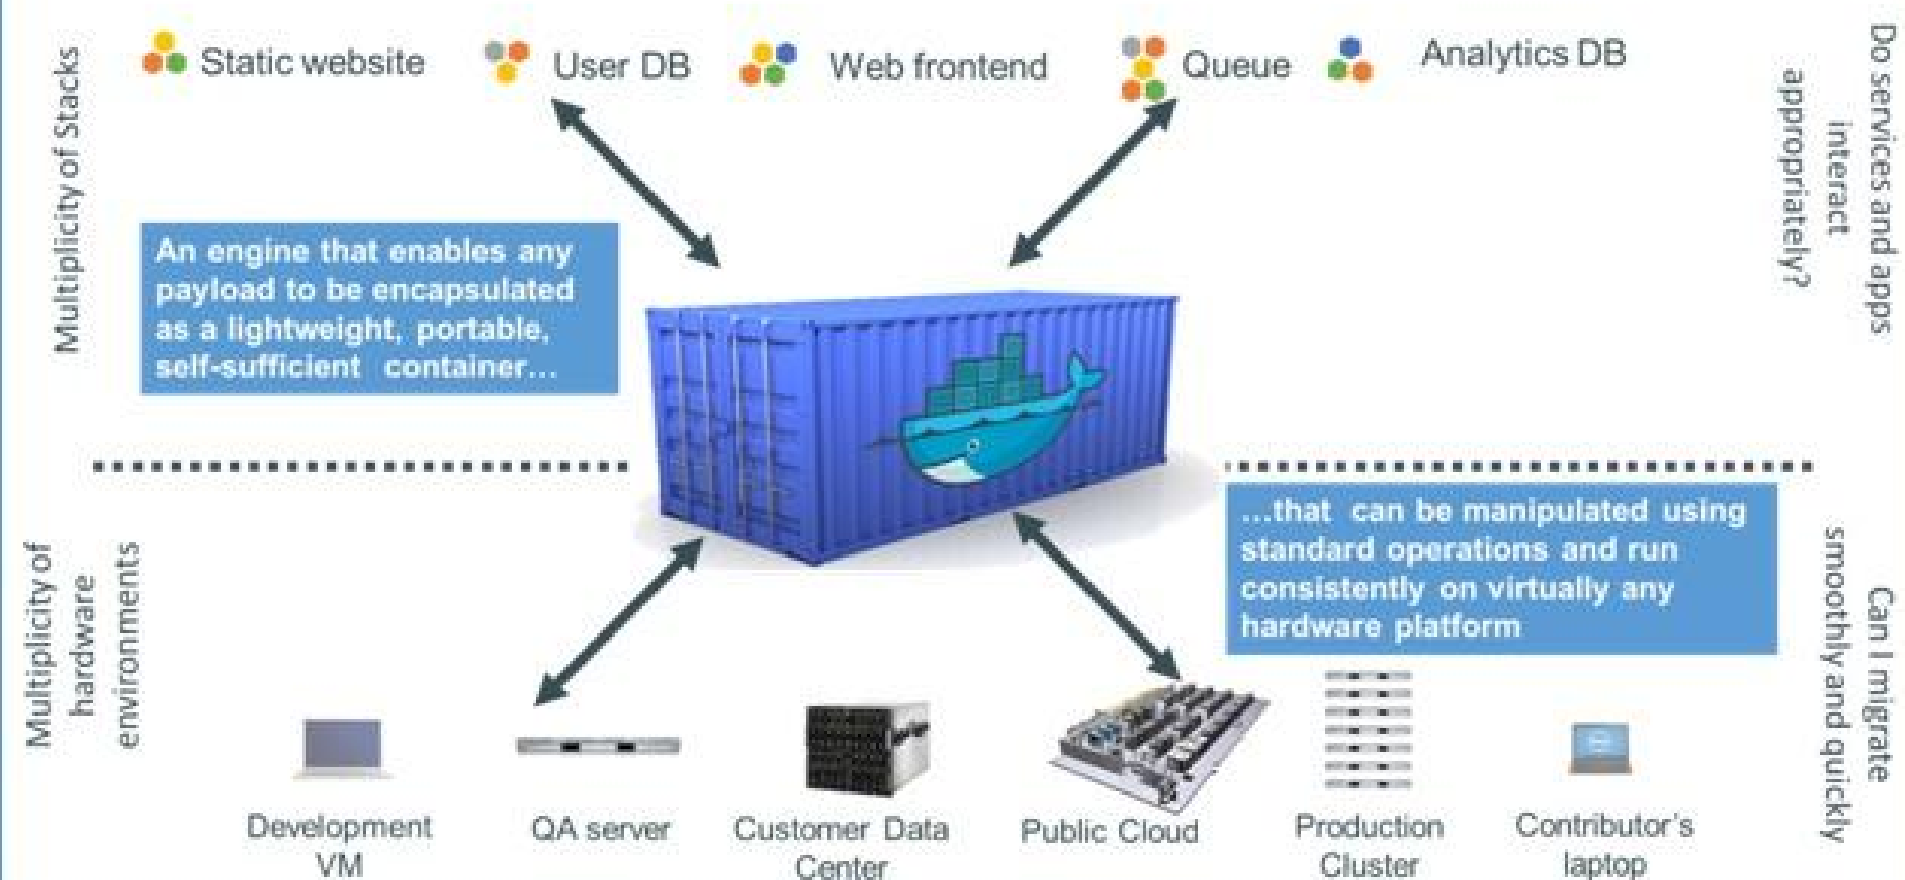
Each virtualized application includes not only the application - which may be only 10s of MB - and the necessary binaries and libraries, but also an entire guest operating system - which may weigh 10s of GB.

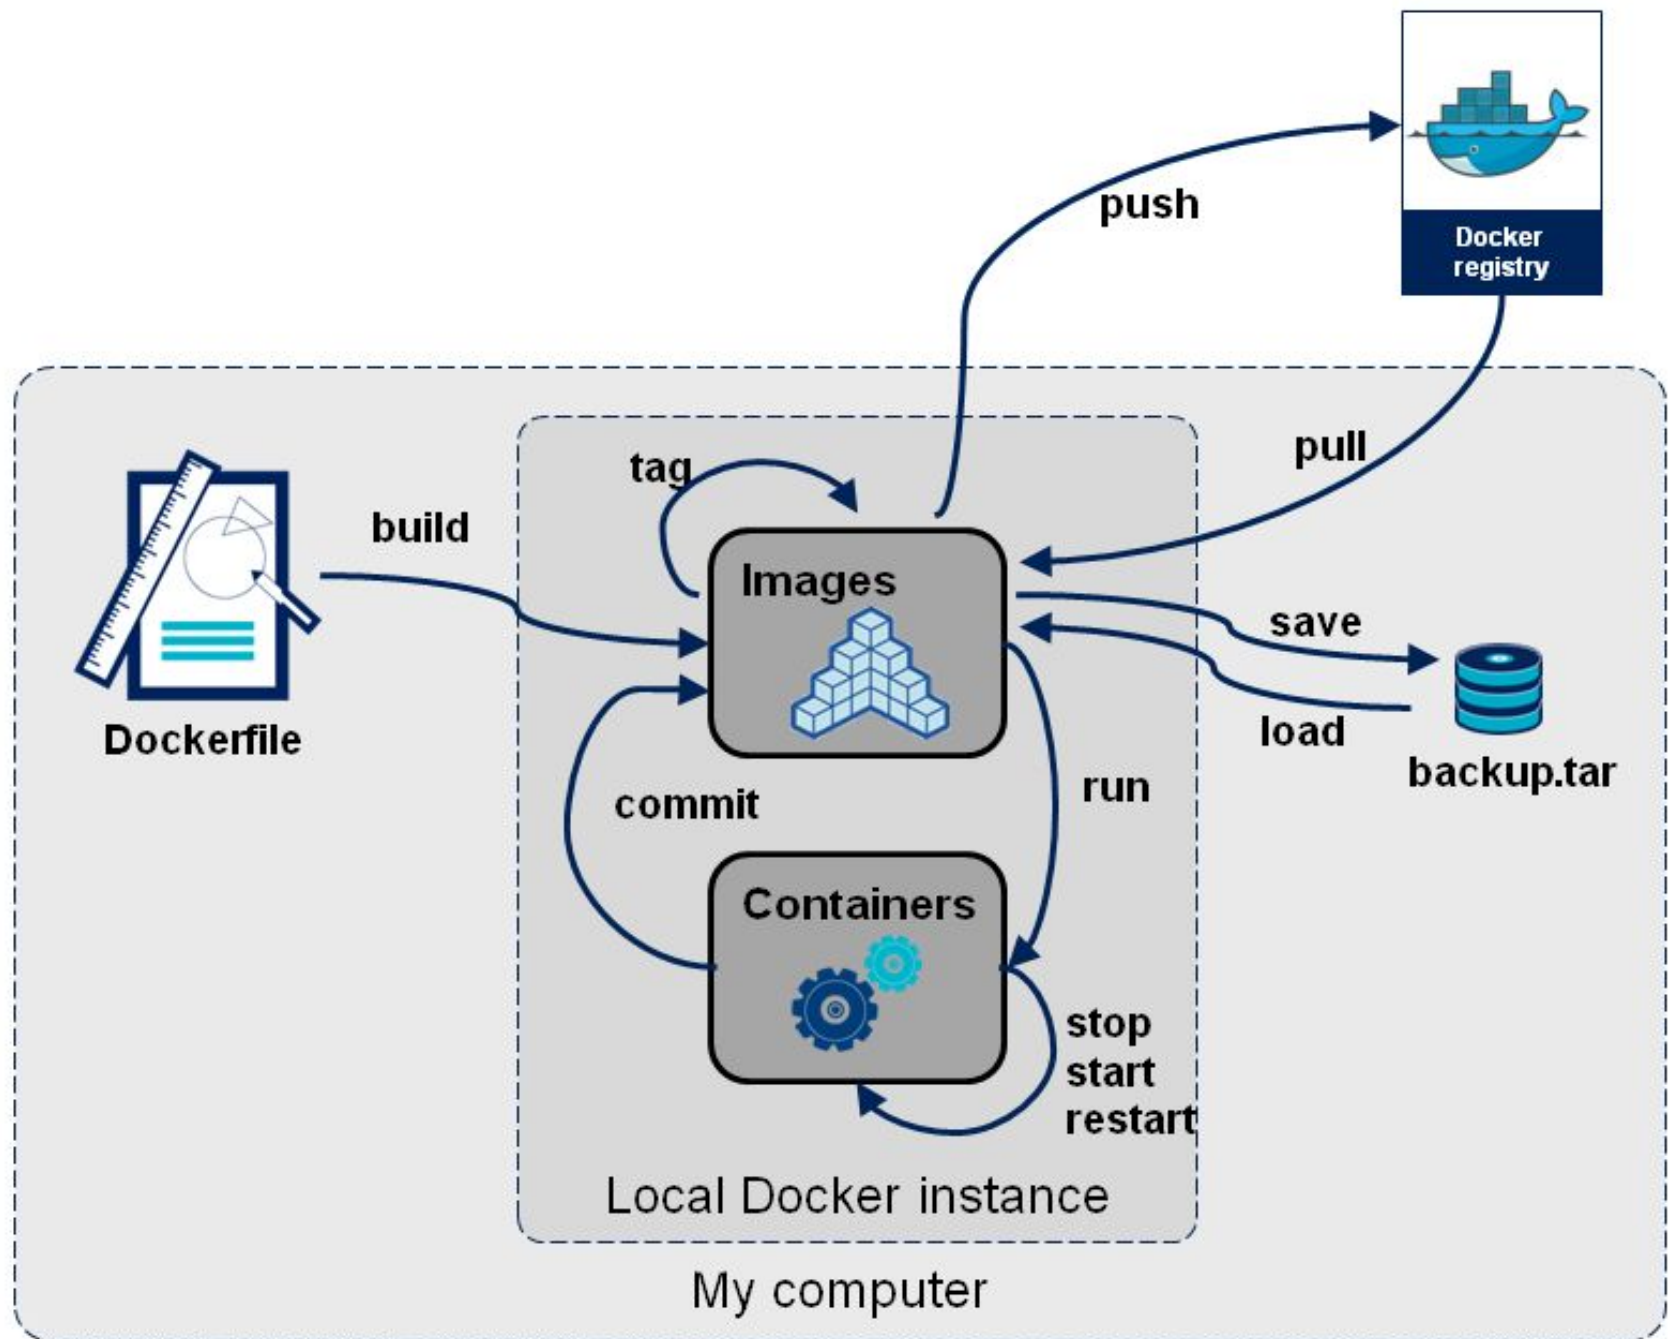


Docker

The Docker Engine container comprises just the application and its dependencies. It runs as an isolated process in userspace on the host operating system, sharing the kernel with other containers. Thus, it enjoys the resource isolation and allocation benefits of VMs but is much more portable and efficient.

Docker is a shipping container system for code





Docker Commands

<https://www.docker.com/tryit/>

<https://docs.docker.com/reference/commandline/cli/>

- a. `docker search centos`
- b. `docker pull centos` (or any other image)
- c. `docker images`
- d. `docker run -it <image name>`
- e. `docker run -v </mymachine>:</vm> <image name> <cmd>`
- f. `docker run -d -p 8080:80 <image name> <cmd>`
- g. `docker ps (-a)`
- h. `docker commit <image id> <repo name>`
- i. `docker build <folder with Docker file>` ([link](#))
- j. `docker save <container id> > out.tar`
- k. `docker import out.tar`
- l. `docker save <container id> | gzip > out.tar.gz`
- m. `docker rm <container id>`

Programming Language-Level Virtualization

- Allows programs to be run on different machines with different OS/architectures.
- Uses virtual machine to 'run' code: slow.
- Java (JVM), .Net (CLI), Python, Pascal, Groovy, Ruby
- Popular with enterprise apps
- Can be interpreted or jitted
- Stack-based (Java, CLI) or register-based
- Security a big win: easy to sandbox

Application-level Virtualization

- Allows applications to run in an environment it shouldn't
- Extra layer emulates missing pieces: filesystem, libraries, OS
- *2 methods:*
 - *Interpretation:* every ISA instruction is handled - minimal startup, huge overhead
 - *Binary translation:* Instructions translated to native, then cached and reused.
- Ex: WINE, CrossOver (Mac), VMware ThinApp packages an exe



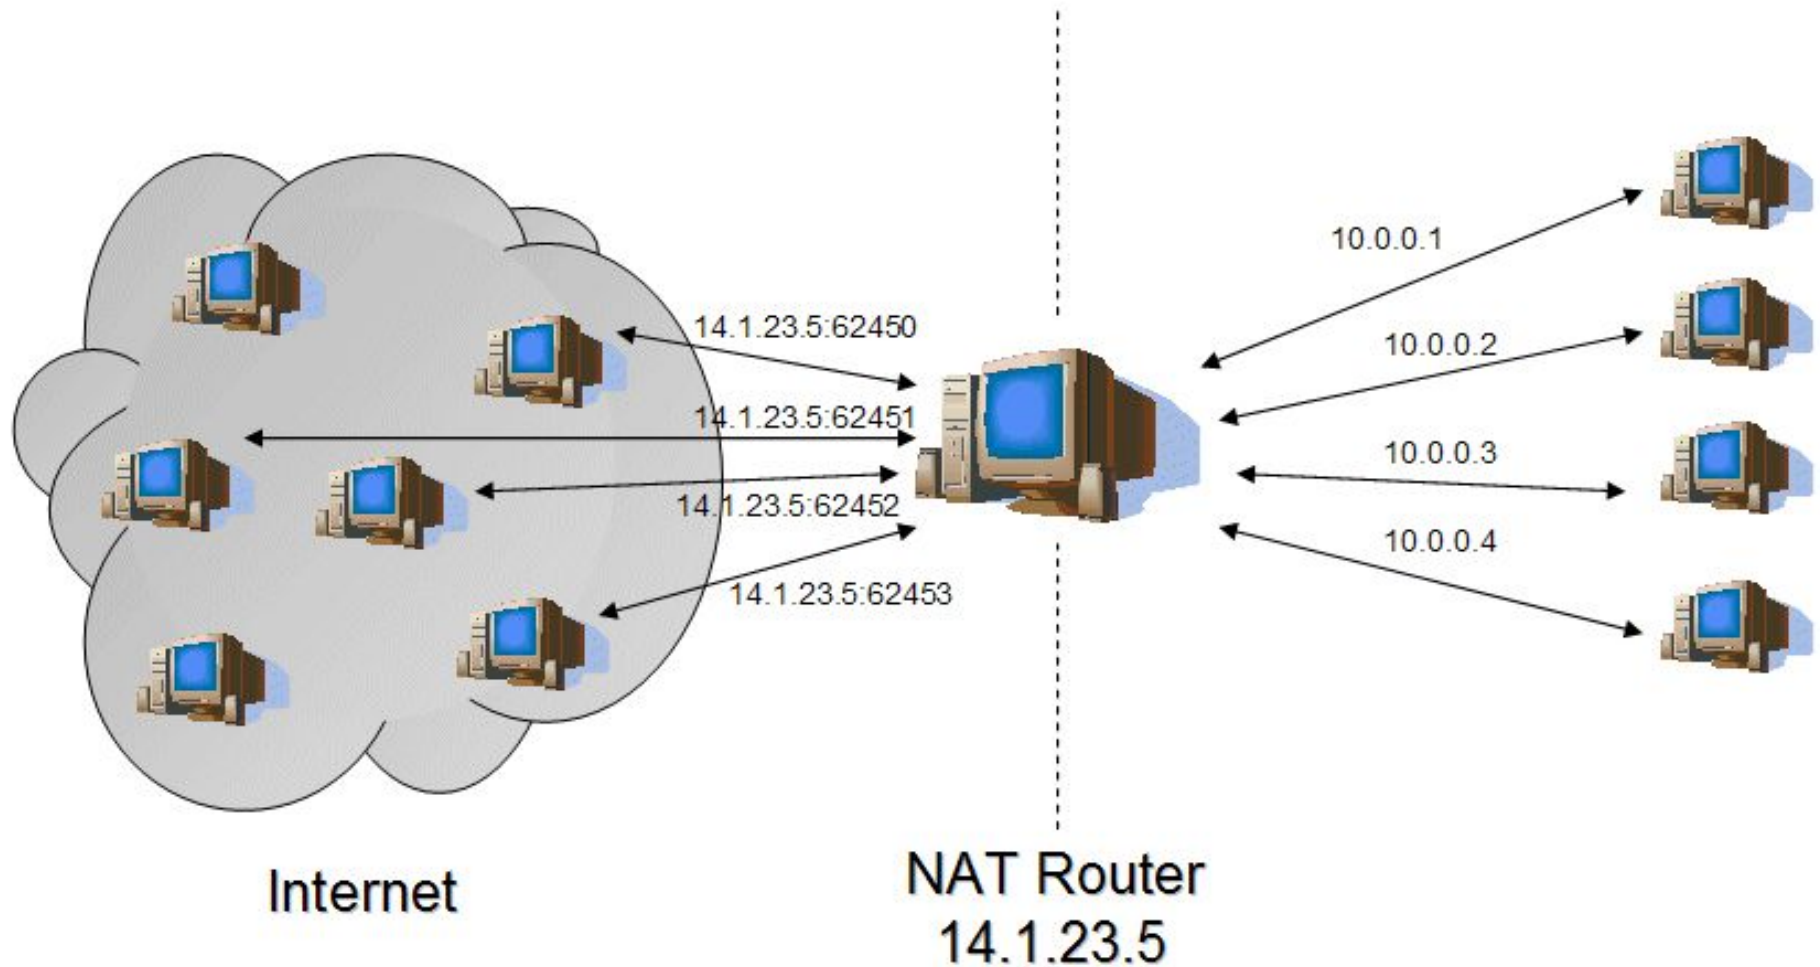
Storage Virtualization

- Physical location of data not important
- Accessed through logical path or similar
- Management of write access important
- Ex:
 - [SAN](#) - Virtualize SATA/SCSI cable - single computer at a time connected
 - [NAS](#) - Expose virtual filesystem over a network (NFS, AFP, SMB)
 - Dropbox.com, Box.com, SkyDrive

Network Virtualization

- *External* net virtualization - VLAN - Separate layer 2 networks (broadcast domain) on same physical wire
- *Internal* net virtualization - simulated network in computer, as in to connect multiple VMs on in one host.
 - NAT (Network Address Translation)

NAT Diagram



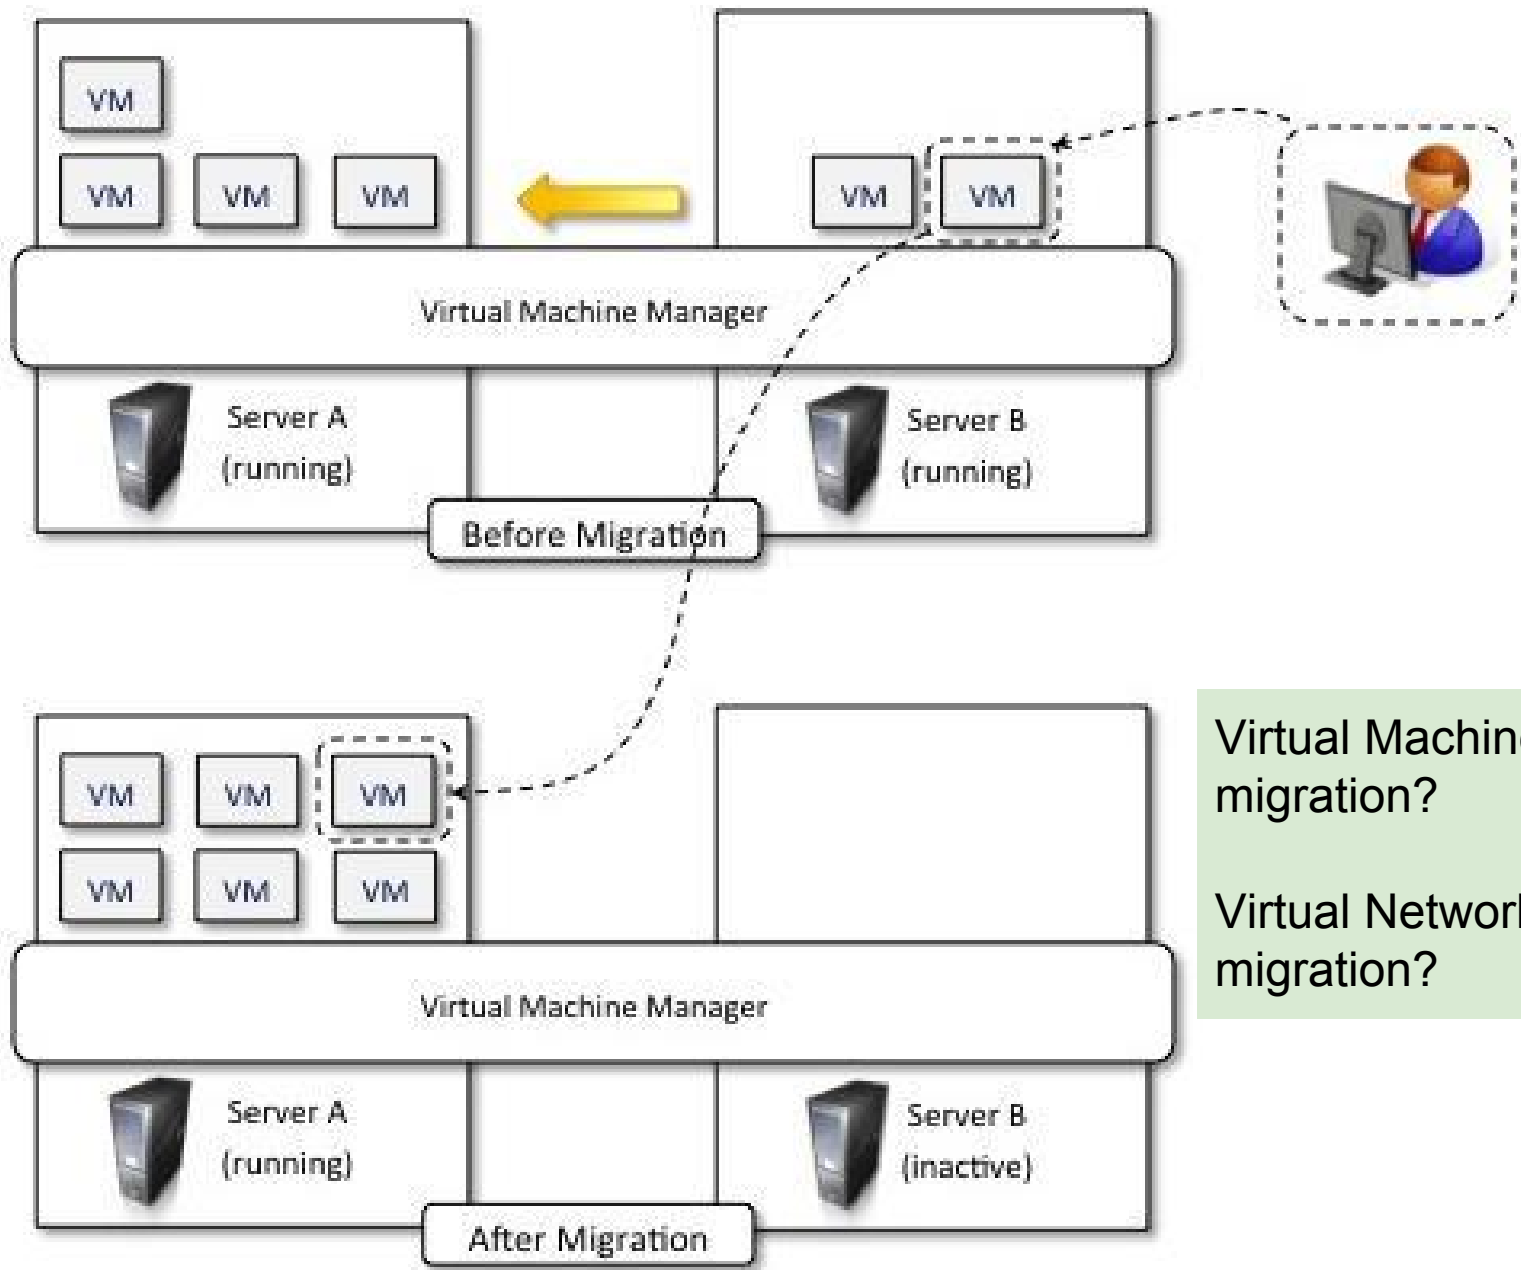
Desktop Virtualization

- Allows a desktop environment remotely
- Pros: high availability, persistence, accessibility, ease of management
- Neg: lower refresh, lag
- Old idea... from mainframe era
- Implemented with IaaS or in single OS
- Ex: RDP (Windows Remote Services), VNC, X Server.

Virtualization and Cloud Computing

- Facilitates IaaS
- Facilitates PaaS
- Customization, security, isolation, manageability
- Allows better use of abstraction
- Opens new market-based computing opportunities

VM Migration



Virtualization Pros/Cons

- Pros:
 - Managed execution
 - Isolation
 - Portability / Migration
 - Efficiency
- Cons:
 - Efficiency - hypervisor or degraded environment
 - Security
 - Phishing
 - Malicious hypervisor: [BluPill](#), SubVirt (rootkits)
 - Replace VM with compromised one (JVM)

Technology Examples - FYI

Not Tested - Details in Text

- Xen: paravirtualization
 - XenSource - Citrix (commercial)
 - Xen Cloud Platform - XCP - open source
- VMware: Full virtualization
 - Desktop
 - VMware Workstation
 - VMware Fusion (Mac)
 - Server
 - VMware GSX
 - VMware ESXi
 - VMware vSphere - cloud management solution

Technology Examples - FYI

Not Tested - Details in Text

- Microsoft Hyper-V
 - Part of Windows Server 2008 R2 onward
 - Windows Server Core (2008) Optimized for Hyper-V (removed extra MS components)
 - System Center Virtual Machine Manager (SCVMM) - manage VM's on multiple servers - can interoperate with VMware

Virtualization Summary

- Ability to provide illusion of an environment
 - Runtime:
 - Hardware: IaaS - VMware, VirtualBox
 - OS Level: CHROOT, Docker
 - Programming language: JVM, Python
 - Application: WINE
 - Storage: SAN, NAS, Dropbox.com
 - Network: VLAN, Virtual Networking, NAT, VPN
 - Remote Desktop: RDP, X Server
 - Pros/cons